

```

# prompt: import data from kaggle

# Install the Kaggle API client.
!pip install kaggle

# Upload your Kaggle API key (kaggle.json).
# Go to https://www.kaggle.com/<username>/account, click "Create New
API Token", and upload the downloaded kaggle.json file.
from google.colab import files
files.upload()

# Move the kaggle.json file to the correct location.
!mkdir -p ~/.kaggle
!mv kaggle.json ~/.kaggle/

# Change the permissions of the file.
!chmod 600 ~/.kaggle/kaggle.json

# Download the dataset. Replace 'username/dataset-name' with the
actual username and dataset name.
!kaggle datasets download -d 'shantanudhakadd/diabetes-dataset-for-
beginners'

# Unzip the downloaded file (if necessary).
!unzip doctors-handwritten-prescription-bd-dataset.zip # replace with
actual filename if different

# Now you can load the data using pandas (or other libraries)
import pandas as pd

# Example: Load a CSV file
# df = pd.read_csv("your_data_file.csv") # Replace with your actual
data file name
# print(df.head())

Requirement already satisfied: kaggle in
/usr/local/lib/python3.10/dist-packages (1.6.17)
Requirement already satisfied: six>=1.10 in
/usr/local/lib/python3.10/dist-packages (from kaggle) (1.16.0)
Requirement already satisfied: certifi>=2023.7.22 in
/usr/local/lib/python3.10/dist-packages (from kaggle) (2024.8.30)
Requirement already satisfied: python-dateutil in
/usr/local/lib/python3.10/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in
/usr/local/lib/python3.10/dist-packages (from kaggle) (2.32.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-
packages (from kaggle) (4.66.6)
Requirement already satisfied: python-slugify in
/usr/local/lib/python3.10/dist-packages (from kaggle) (8.0.4)
Requirement already satisfied: urllib3 in

```

```

/usr/local/lib/python3.10/dist-packages (from kaggle) (2.2.3)
Requirement already satisfied: bleach in
/usr/local/lib/python3.10/dist-packages (from kaggle) (6.2.0)
Requirement already satisfied: webencodings in
/usr/local/lib/python3.10/dist-packages (from bleach->kaggle) (0.5.1)
Requirement already satisfied: text-unidecode>=1.3 in
/usr/local/lib/python3.10/dist-packages (from python-slugify->kaggle)
(1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests->kaggle)
(3.4.0)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.10/dist-packages (from requests->kaggle) (3.10)
<IPython.core.display.HTML object>

Saving kaggle.json to kaggle.json
Dataset URL: https://www.kaggle.com/datasets/shantanudhakadd/diabetes-
dataset-for-beginners
License(s): other
Downloading diabetes-dataset-for-beginners.zip to /content
 0% 0.00/8.91k [00:00<?, ?B/s]
100% 8.91k/8.91k [00:00<00:00, 16.7MB/s]
unzip: cannot find or open doctors-handwritten-prescription-bd-
dataset.zip, doctors-handwritten-prescription-bd-dataset.zip.zip or
doctors-handwritten-prescription-bd-dataset.zip.ZIP.

```

## Read Data

```

data_file = '/content/diabetes-dataset-for-beginners.zip'
data = pd.read_csv(data_file)

# Print column names to verify
print(data.columns)

Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
       'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
import seaborn as sns

data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767

```

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

dtypes: float64(2), int64(7)

memory usage: 54.1 KB

data.head()

```
{
  "summary": {
    "name": "data",
    "rows": 768,
    "fields": [
      {
        "column": "Pregnancies",
        "properties": {
          "dtype": "number",
          "std": 3,
          "min": 0,
          "max": 17,
          "num_unique_values": 17,
          "samples": [6, 1, 3],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "Glucose",
        "properties": {
          "dtype": "number",
          "std": 31,
          "min": 0,
          "max": 199,
          "num_unique_values": 136,
          "samples": [151, 101, 112],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "BloodPressure",
        "properties": {
          "dtype": "number",
          "std": 19,
          "min": 0,
          "max": 122,
          "num_unique_values": 47,
          "samples": [86, 46, 85],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "SkinThickness",
        "properties": {
          "dtype": "number",
          "std": 15,
          "min": 0,
          "max": 99,
          "num_unique_values": 51,
          "samples": [7, 12, 48],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "Insulin",
        "properties": {
          "dtype": "number",
          "std": 115,
          "min": 0,
          "max": 846,
          "num_unique_values": 186,
          "samples": [52, 41, 183],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "BMI",
        "properties": {
          "dtype": "number",
          "std": 7.8841603203754405,
          "min": 0.0,
          "max": 67.1,
          "num_unique_values": 248,
          "samples": [31.0, 38.1, 19.9],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "DiabetesPedigreeFunction",
        "properties": {
          "dtype": "float64",
          "std": 0.47147,
          "min": 0.0,
          "max": 2.674,
          "num_unique_values": 137,
          "samples": [1.0, 1.6, 0.171],
          "semantic_type": "float64",
          "description": ""
        }
      },
      {
        "column": "Age",
        "properties": {
          "dtype": "int64",
          "std": 11.959,
          "min": 21,
          "max": 81,
          "num_unique_values": 61,
          "samples": [39, 27, 54],
          "semantic_type": "int64",
          "description": ""
        }
      },
      {
        "column": "Outcome",
        "properties": {
          "dtype": "int64",
          "std": 0.371,
          "min": 0,
          "max": 1,
          "num_unique_values": 2,
          "samples": [150, 618],
          "semantic_type": "int64",
          "description": ""
        }
      }
    ]
  }
}
```

```

\"semantic_type\": \"\",
  {
    \"column\": \"DiabetesPedigreeFunction\",
    \"properties\": {
      \"dtype\": \"number\",
      \"std\": 0.33132859501277484,
      \"min\": 0.078,
      \"max\": 2.42,
      \"num_unique_values\": 517,
      \"samples\": [
        1.731,
        0.426,
        0.138
      ]
    }
  },
  {
    \"column\": \"Age\",
    \"properties\": {
      \"dtype\": \"number\",
      \"std\": 11,
      \"min\": 21,
      \"max\": 81,
      \"num_unique_values\": 52,
      \"samples\": [
        60,
        47,
        72
      ]
    }
  },
  {
    \"column\": \"Outcome\",
    \"properties\": {
      \"dtype\": \"number\",
      \"std\": 0,
      \"min\": 0,
      \"max\": 1,
      \"num_unique_values\": 2,
      \"samples\": [
        0,
        1
      ]
    }
  }
],
\"type\": \"dataframe\",
\"variable_name\": \"data\"}

```

```
data.describe()
```

```

{
  \"summary\": {
    \"name\": \"data\",
    \"rows\": 8,
    \"fields\": [
      {
        \"column\": \"Pregnancies\",
        \"properties\": {
          \"dtype\": \"number\",
          \"std\": 269.85223453356366,
          \"min\": 0.0,
          \"max\": 768.0,
          \"num_unique_values\": 8,
          \"samples\": [
            3.8450520833333335,
            3.0,
            768.0
          ]
        },
        \"semantic_type\": \"\",
        \"description\": \"\"
      },
      {
        \"column\": \"Glucose\",
        \"properties\": {
          \"dtype\": \"number\",
          \"std\": 243.73802348295857,
          \"min\": 0.0,
          \"max\": 768.0,
          \"num_unique_values\": 8,
          \"samples\": [
            120.89453125,
            117.0,
            768.0
          ]
        },
        \"semantic_type\": \"\",
        \"description\": \"\"
      },
      {
        \"column\": \"BloodPressure\",
        \"properties\": {
          \"dtype\": \"number\",
          \"std\": 252.85250535810619,
          \"min\": 0.0,
          \"max\": 768.0,
          \"num_unique_values\": 8,
          \"samples\": [
            69.10546875,
            72.0,
            768.0
          ]
        },
        \"semantic_type\": \"\",
        \"description\": \"\"
      },
      {
        \"column\": \"SkinThickness\",
        \"properties\": {
          \"dtype\": \"number\",
          \"std\": 263.7684730531098,
          \"min\": 0.0,
          \"max\": 768.0,
          \"num_unique_values\": 7,
          \"samples\": [
            768.0,
            20.536458333333332,
            32.0
          ]
        },
        \"semantic_type\": \"\",
        \"description\": \"\"
      },
      {
        \"column\": \"Insulin\",
        \"properties\": {
          \"dtype\": \"number\",
          \"std\": 350.26059167945886,
          \"min\": 0.0,
          \"max\": 846.0,
          \"num_unique_values\": 7,
          \"samples\": [
            768.0,
            79.79947916666667,
            79.79947916666667
          ]
        },
        \"semantic_type\": \"\",
        \"description\": \"\"
      }
    ]
  }
}

```

```

127.25\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n        {\n          \"column\": \"BMI\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 262.05117817552093,\n            \"min\": 0.0,\n            \"max\": 768.0,\n            \"num_unique_values\": 8,\n            \"samples\": [\n              31.992578124999998,\n              32.0,\n              768.0\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          },\n          {\n            \"column\": \"DiabetesPedigreeFunction\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 271.3005221658502,\n              \"min\": 0.078,\n              \"max\": 768.0,\n              \"num_unique_values\": 8,\n              \"samples\": [\n                0.47187630208333325,\n                0.3725,\n                768.0\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            },\n            {\n              \"column\": \"Age\",\n              \"properties\": {\n                \"dtype\": \"number\",\n                \"std\": 260.1941178528413,\n                \"min\": 11.76023154067868,\n                \"max\": 768.0,\n                \"num_unique_values\": 8,\n                \"samples\": [\n                  33.240885416666664,\n                  29.0,\n                  768.0\n                ],\n                \"semantic_type\": \"\",\n                \"description\": \"\"\n              },\n              {\n                \"column\": \"Outcome\",\n                \"properties\": {\n                  \"dtype\": \"number\",\n                  \"std\": 271.3865920388932,\n                  \"min\": 0.0,\n                  \"max\": 768.0,\n                  \"num_unique_values\": 5,\n                  \"samples\": [\n                    0.3489583333333333,\n                    1.0,\n                    0.4769513772427971\n                  ],\n                  \"semantic_type\": \"\",\n                  \"description\": \"\"\n                }\n              }\n            }\n          }\n        ],\n        \"type\": \"dataframe\"}

```

```
data.isnull().sum()
```

```

Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64

```

```
# Import Libraries
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
# 1. Correlation Heatmap
```

```
plt.figure(figsize=(10, 8))
```

```
correlation_matrix = data.corr() # Compute correlation matrix
```

```
sns.heatmap(
```

```
    correlation_matrix,
```

```
    annot=True, # Show correlation coefficients
```

```


    fmt=".2f",    # Format float values
    cmap="coolwarm",
    cbar=True
)
plt.title("Correlation Heatmap")
plt.show()

# 2. Pair Plot (Relationships Between Features)
# Only plot selected features for simplicity
selected_features = ['Glucose', 'BMI', 'Age', 'Insulin',
                     'SkinThickness', 'Outcome']
sns.pairplot(data[selected_features], hue='Outcome', diag_kind='kde',
             palette='husl')
plt.show()

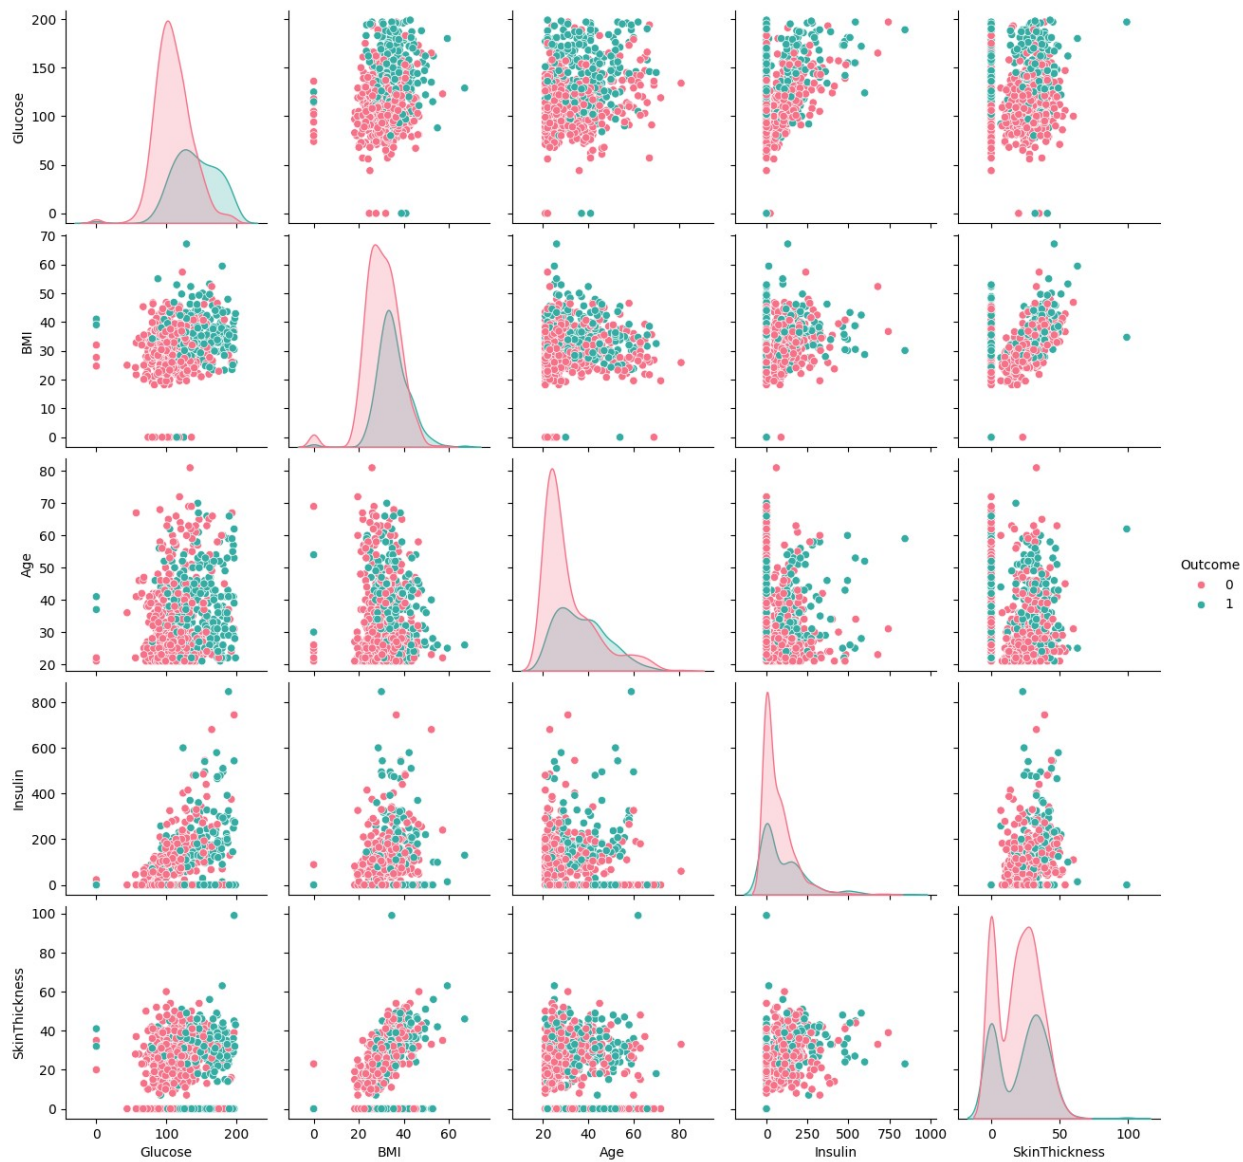
# 3. Feature Distribution Plot
# Visualize distributions of selected features
for feature in selected_features:
    if feature != 'Outcome':
        plt.figure(figsize=(6, 4))
        sns.histplot(data, x=feature, hue="Outcome", kde=True,
                    palette="husl", bins=20)
        plt.title(f"Distribution of {feature} by Outcome")
        plt.xlabel(feature)
        plt.ylabel("Frequency")
        plt.show()

# 4. Box Plot (Feature vs Outcome)
# Analyze the distribution of numeric features based on the target
variable
for feature in selected_features:
    if feature != 'Outcome':
        plt.figure(figsize=(6, 4))
        sns.boxplot(data=data, x='Outcome', y=feature, palette="husl")
        plt.title(f"Box Plot of {feature} vs Outcome")
        plt.xlabel("Outcome")
        plt.ylabel(feature)
        plt.show()

```












Distribution of Glucose by Outcome




Distribution of BMI by Outcome




Distribution of Age by Outcome



Distribution of Insulin by Outcome






```
<ipython-input-7-568432ea0180>:40: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```

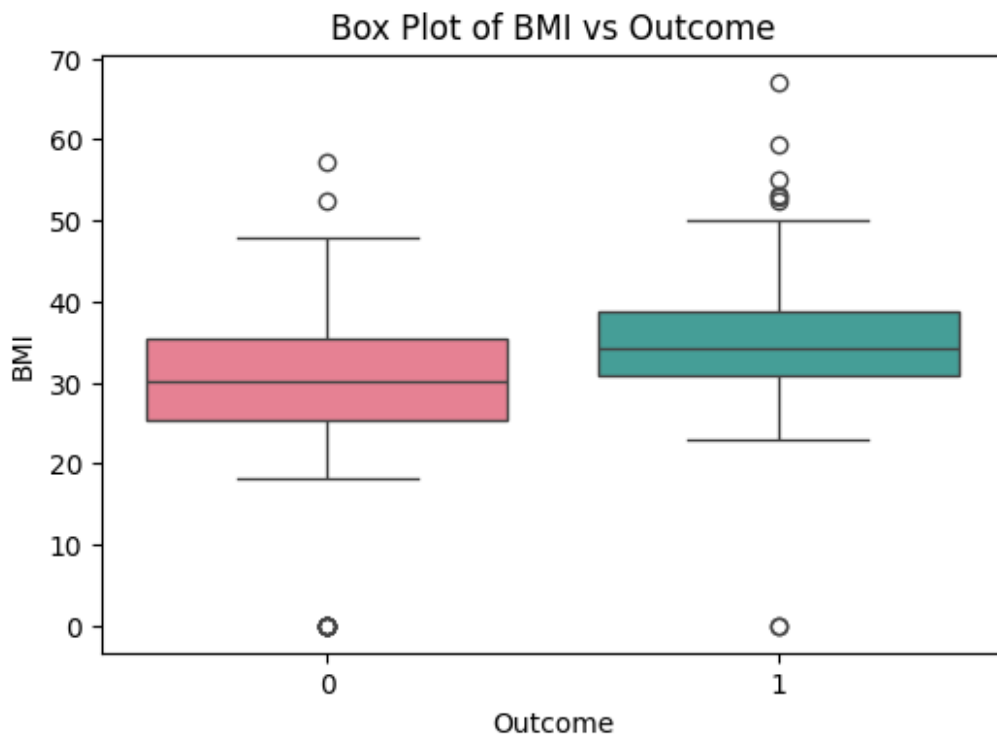
```
sns.boxplot(data=data, x='Outcome', y=feature, palette="husl")
```



```
<ipython-input-7-568432ea0180>:40: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```

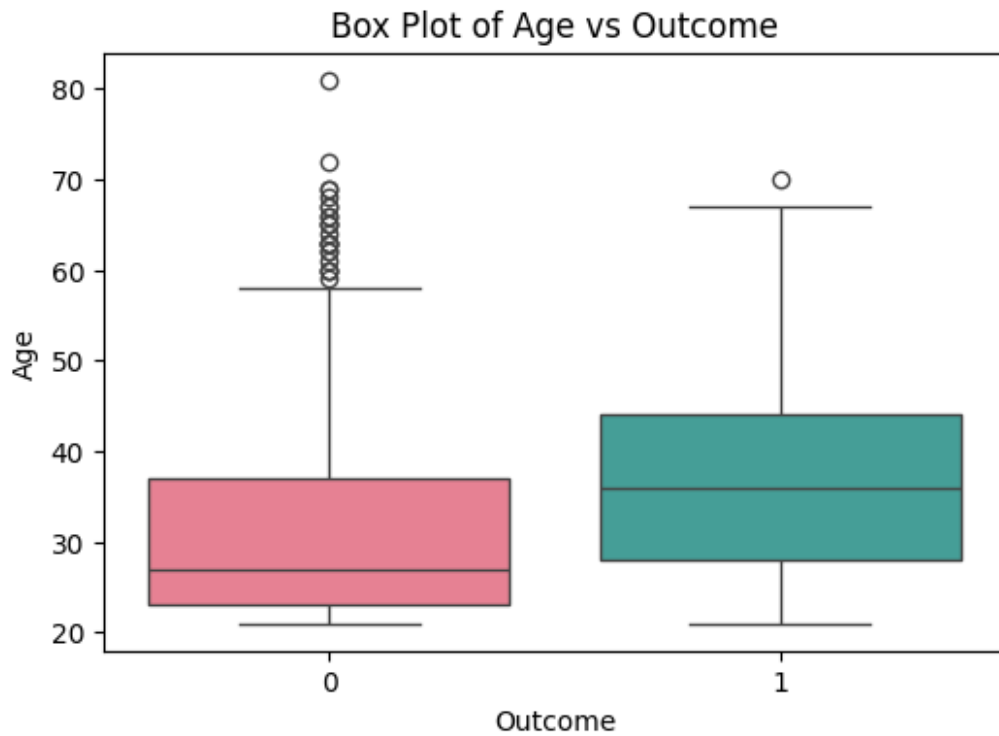
```
sns.boxplot(data=data, x='Outcome', y=feature, palette="husl")
```



```
<ipython-input-7-568432ea0180>:40: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```


```
sns.boxplot(data=data, x='Outcome', y=feature, palette="husl")
```



```
<ipython-input-7-568432ea0180>:40: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```

```
sns.boxplot(data=data, x='Outcome', y=feature, palette="husl")
```




```
<ipython-input-7-568432ea0180>:40: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```

```
sns.boxplot(data=data, x='Outcome', y=feature, palette="husl")
```





```
# let's look at the Corralation matrix of this data
corr_matrix = data.corr()
corr_matrix['Outcome'].sort_values(ascending=False)

Outcome          1.000000
Glucose           0.466581
BMI               0.292695
Age              0.238356
Pregnancies       0.221898
DiabetesPedigreeFunction 0.173844
Insulin           0.130548
SkinThickness     0.074752
BloodPressure     0.065068
Name: Outcome, dtype: float64

X = data.drop('Outcome', axis=1) # Features
y = data['Outcome'] # Target

from sklearn.feature_selection import SelectKBest, f_classif
#Feature Selection (Select Top 5 Features using ANOVA)
selector = SelectKBest(score_func=f_classif, k=5)
X_selected = selector.fit_transform(X, y)
selected_features = X.columns[selector.get_support()]
print("Selected Features:", selected_features.tolist())

Selected Features: ['Pregnancies', 'Glucose', 'BMI',
'DiabetesPedigreeFunction', 'Age']
```

```

#Split the data into training and test sets:
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

#select machine learning model
from sklearn.ensemble import RandomForestClassifier

model = RandomForestClassifier()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

#Model Evaluation

# cross validation
from sklearn.model_selection import cross_val_score
scores = cross_val_score(model, X, y, cv=5)
print("Cross-Validation Accuracy:", scores.mean())

# ROC-AUC curve
from sklearn.metrics import roc_auc_score, roc_curve
y_pred_proba = model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)


plt.plot(fpr, tpr, label='ROC Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC-AUC Curve')
plt.show()

# Confusion Matrix
from sklearn.metrics import confusion_matrix, classification_report
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))

# Precision and Recall
from sklearn.metrics import precision_score, recall_score
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))

Cross-Validation Accuracy: 0.7683388506917919

```



```
[[81 18]
 [18 37]]
```


	precision	recall	f1-score	support
0	0.82	0.82	0.82	99
1	0.67	0.67	0.67	55
accuracy			0.77	154
macro avg	0.75	0.75	0.75	154
weighted avg	0.77	0.77	0.77	154

Precision: 0.6727272727272727

Recall: 0.6727272727272727

*# Visualize important features using the model's feature importances (e.g., for Random Forest)*

```
feature_importances = model.feature_importances_
sns.barplot(x=feature_importances, y=X.columns)
plt.show()
```



```
# Save the model using joblib for deployment:
import joblib
joblib.dump(model, 'random_forest_model.pkl')

# Load the saved model
loaded_model = joblib.load('/content/random_forest_model.pkl')

# 1. Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import (
    SelectKBest,
    f_classif, # ANOVA F-value
    RFE
)
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression, Lasso
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, roc_auc_score,
roc_curve, confusion_matrix
from xgboost import XGBClassifier

# Advanced Feature Selection Function
def advanced_feature_selection(X, y):
```

```

# Methods to compare
feature_selection_methods = {
    'Univariate Selection (ANOVA)':
SelectKBest(score_func=f_classif, k=5),
    'Recursive Feature Elimination': RFE(
        estimator=RandomForestClassifier(),
        n_features_to_select=5
    ),
    'L1 Regularization (Lasso)': Lasso(alpha=0.1)
}

# Store selected features
selected_features = {}
feature_importances = {}

# Standardize features for fair comparison
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Perform feature selection
for name, selector in feature_selection_methods.items():
    # Fit selector
    if name == 'L1 Regularization (Lasso)':
        selector.fit(X_scaled, y)
        # Get feature importances for Lasso
        importances = np.abs(selector.coef_)
        selected_mask = importances > 0
        selected_features[name] = X.columns[selected_mask]
        feature_importances[name] = pd.Series(
            importances[selected_mask],
            index=X.columns[selected_mask]
        )
    elif hasattr(selector, 'get_support'):
        selector.fit(X_scaled, y)
        selected_features[name] =
X.columns[selector.get_support()]

        # For RFE, get feature ranking
        if name == 'Recursive Feature Elimination':
            feature_importances[name] = pd.Series(
                selector.ranking_,
                index=X.columns
            )
    else:
        selector.fit(X_scaled, y)
        selected_features[name] = X.columns

# Visualize feature importances
plt.figure(figsize=(12, 6))
for i, (name, importances) in

```

```

enumerate(feature_importances.items(), 1):
    plt.subplot(1, len(feature_importances), i)
    importances.sort_values(ascending=False).plot(kind='bar')
    plt.title(f'Feature Importance: {name}')
    plt.xticks(rotation=45, ha='right')
    plt.tight_layout()
plt.show()

return selected_features

# 2. Load Dataset
data = pd.read_csv("/content/diabetes-dataset-for-beginners.zip")

# Check for missing values
print("Missing Values:\n", data.isnull().sum())

# Fill missing values with mean
data.fillna(data.mean(), inplace=True)

# Encode categorical variables (if any)
data = pd.get_dummies(data, drop_first=True)

# Perform advanced feature selection
X = data.drop('Outcome', axis=1)
y = data['Outcome']

# Get selected features from different methods
selected_feature_sets = advanced_feature_selection(X, y)

# Print selected features
print("\nSelected Features:")
for method, features in selected_feature_sets.items():
    print(f"\n{method}:")
    print(features.tolist())

# Use the intersection of features from different methods
common_features = list(set.intersection(
    *[set(features) for features in selected_feature_sets.values()]
))

# If no common features, use features from the best method
if not common_features:
    # Prioritize methods in this order
    priority_methods = [
        'L1 Regularization (Lasso)',
        'Recursive Feature Elimination',
        'Univariate Selection (ANOVA)'
    ]

    for method in priority_methods:

```

```

        if len(selected_feature_sets[method]) > 0:
            common_features = list(selected_feature_sets[method])
            break

# Update X with selected features
X_selected = X[common_features]

# 3. Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X_selected, y,
test_size=0.2, random_state=42)

# 4. Define Models
models = {
    "Logistic Regression": LogisticRegression(),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "SVM": SVC(probability=True),
    "XGBoost": XGBClassifier(eval_metric='logloss')
}

# 5. Define Pipeline
pipelines = {
    name: Pipeline([
        ('scaler', StandardScaler()),
        ('classifier', model)
    ])
    for name, model in models.items()
}

# 6. Train and Evaluate Models
results = {}
for name, pipeline in pipelines.items():
    print(f"\nTraining {name}...")
    pipeline.fit(X_train, y_train)
    y_pred = pipeline.predict(X_test)
    y_pred_proba = pipeline.predict_proba(X_test)[: , 1] if
hasattr(pipeline, 'predict_proba') else None

    # Evaluation Metrics
    print(f"Classification Report for {name}:\n",
classification_report(y_test, y_pred))
    confusion = confusion_matrix(y_test, y_pred)
    print(f"Confusion Matrix for {name}:\n", confusion)

    # ROC-AUC Score
    if y_pred_proba is not None:
        roc_auc = roc_auc_score(y_test, y_pred_proba)
        print(f"ROC-AUC Score for {name}: {roc_auc:.2f}")
        results[name] = {'roc_auc': roc_auc}

```



```

    # Plot ROC Curve
    fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
    plt.plot(fpr, tpr, label=f"{name} (AUC = {roc_auc:.2f})")
else:
    results[name] = {'roc_auc': None}

```

```

# Show ROC Curves

```

```

plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.title("ROC Curves")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.show()

```

```

# 7. Model Comparison

```

```

comparison = pd.DataFrame(results).T.sort_values(by='roc_auc',
ascending=False)
print("\nModel Performance Comparison:\n", comparison)

```

```

# 8. Save the Best Model

```

```

best_model_name = comparison.index[0]
best_pipeline = pipelines[best_model_name]
import joblib
joblib.dump(best_pipeline, f"{best_model_name}_pipeline.pkl")
print(f"\nBest Model ({best_model_name}) saved as
'{best_model_name}_pipeline.pkl'.")

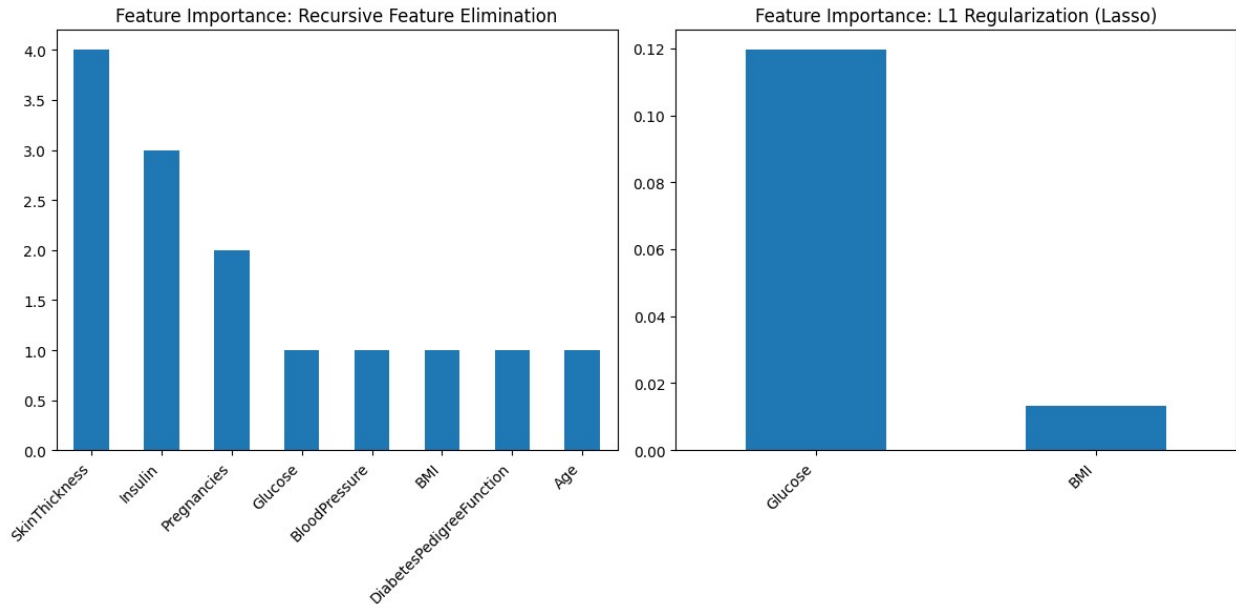
```

```

Missing Values:

```

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0
dtype:	int64



Selected Features:

Univariate Selection (ANOVA):

['Pregnancies', 'Glucose', 'BMI', 'DiabetesPedigreeFunction', 'Age']

Recursive Feature Elimination:

['Glucose', 'BloodPressure', 'BMI', 'DiabetesPedigreeFunction', 'Age']

L1 Regularization (Lasso):

['Glucose', 'BMI']

Training Logistic Regression...

Classification Report for Logistic Regression:

	precision	recall	f1-score	support
0	0.80	0.85	0.82	99
1	0.69	0.62	0.65	55
accuracy			0.77	154
macro avg	0.75	0.73	0.74	154
weighted avg	0.76	0.77	0.76	154

Confusion Matrix for Logistic Regression:

[[84 15]

[21 34]]

ROC-AUC Score for Logistic Regression: 0.81

Training Decision Tree...

Classification Report for Decision Tree:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.78	0.81	0.79	99
1	0.63	0.58	0.60	55
accuracy			0.73	154
macro avg	0.70	0.69	0.70	154
weighted avg	0.72	0.73	0.72	154

Confusion Matrix for Decision Tree:

```
[[80 19]
```

```
[23 32]]
```

ROC-AUC Score for Decision Tree: 0.69

Training Random Forest...

Classification Report for Random Forest:

	precision	recall	f1-score	support
0	0.79	0.83	0.81	99
1	0.66	0.60	0.63	55
accuracy			0.75	154
macro avg	0.72	0.71	0.72	154
weighted avg	0.74	0.75	0.74	154

Confusion Matrix for Random Forest:

```
[[82 17]
```

```
[22 33]]
```

ROC-AUC Score for Random Forest: 0.80

Training SVM...

Classification Report for SVM:

	precision	recall	f1-score	support
0	0.79	0.86	0.83	99
1	0.70	0.60	0.65	55
accuracy			0.77	154
macro avg	0.75	0.73	0.74	154
weighted avg	0.76	0.77	0.76	154

Confusion Matrix for SVM:

```
[[85 14]
```

```
[22 33]]
```


ROC-AUC Score for SVM: 0.79

Training XGBoost...

Classification Report for XGBoost:

	precision	recall	f1-score	support
0	0.79	0.79	0.79	99

	1	0.62	0.62	0.62	55
accuracy				0.73	154
macro avg		0.70	0.70	0.70	154
weighted avg		0.73	0.73	0.73	154
Confusion Matrix for XGBoost:					
[[78 21]					
[21 34]]					
ROC-AUC Score for XGBoost: 0.78					



Model Performance Comparison:	
	roc_auc
Logistic Regression	0.808815
Random Forest	0.799541
SVM	0.793205
XGBoost	0.779706
Decision Tree	0.694949
Best Model (Logistic Regression) saved as 'Logistic Regression_pipeline.pkl'.	

```

# 1. Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, roc_auc_score,
roc_curve, confusion_matrix
from xgboost import XGBClassifier
import joblib

# 2. Load and Preprocess Data
data = pd.read_csv("/content/diabetes-dataset-for-beginners.zip") #
Replace with your dataset path
data.fillna(data.mean(), inplace=True) # Fill missing values
data = pd.get_dummies(data, drop_first=True) # Encode categorical
variables

X = data.drop('Outcome', axis=1) # Features
y = data['Outcome'] # Target

# 3. Feature Selection (Select Top 5 Features using ANOVA)
selector = SelectKBest(score_func=f_classif, k=5)
X_selected = selector.fit_transform(X, y)
selected_features = X.columns[selector.get_support()]
print("Selected Features:", selected_features.tolist())

# 4. Train-Test Split
X_train, X_test, y_train, y_test =
train_test_split(X[selected_features], y, test_size=0.2,
random_state=42)

# 5. Define Models
models = {
    "Logistic Regression": LogisticRegression(),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "SVM": SVC(probability=True), # Enable probability for ROC
    "XGBoost": XGBClassifier(eval_metric='logloss')
}

# 6. Train and Evaluate Models
results = {}
for name, model in models.items():

```

```

print(f"\nTraining {name}...")

# Create pipeline: Scaling -> Model
pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('classifier', model)
])
pipeline.fit(X_train, y_train) # Train

# Predictions and Probabilities
y_pred = pipeline.predict(X_test)
y_pred_proba = pipeline.predict_proba(X_test)[:, 1] if
hasattr(pipeline, 'predict_proba') else None

# Metrics
print(f"\n{name} Classification Report:\n",
classification_report(y_test, y_pred))
print(f"{name} Confusion Matrix:\n", confusion_matrix(y_test,
y_pred))

# ROC-AUC
if y_pred_proba is not None:
    roc_auc = roc_auc_score(y_test, y_pred_proba)
    print(f"{name} ROC-AUC Score: {roc_auc:.2f}")
    results[name] = roc_auc

# Plot ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
plt.plot(fpr, tpr, label=f"{name} (AUC = {roc_auc:.2f})")

# Plot all ROC Curves
plt.plot([0, 1], [0, 1], 'r--')
plt.title("ROC Curves")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.show()

# 7. Compare and Save Best Model
best_model_name = max(results, key=results.get) # Get model with
highest ROC-AUC
print(f"\nBest Model: {best_model_name} with AUC =
{results[best_model_name]:.2f}")

# Save the best model pipeline
best_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('classifier', models[best_model_name])
])
best_pipeline.fit(X_train, y_train)

```

```
joblib.dump(best_pipeline, f"{best_model_name}_pipeline.pkl")
print(f"Best model saved as {best_model_name}_pipeline.pkl")
```

Selected Features: ['Pregnancies', 'Glucose', 'BMI',  
'DiabetesPedigreeFunction', 'Age']

Training Logistic Regression...

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.80	0.82	0.81	99
1	0.66	0.64	0.65	55
accuracy			0.75	154
macro avg	0.73	0.73	0.73	154
weighted avg	0.75	0.75	0.75	154

Logistic Regression Confusion Matrix:

```
[[81 18]
```

```
[20 35]]
```

Logistic Regression ROC-AUC Score: 0.81

Training Decision Tree...

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	0.80	0.74	0.77	99
1	0.59	0.67	0.63	55
accuracy			0.71	154
macro avg	0.69	0.71	0.70	154
weighted avg	0.73	0.71	0.72	154

Decision Tree Confusion Matrix:

```
[[73 26]
```

```
[18 37]]
```

Decision Tree ROC-AUC Score: 0.71

Training Random Forest...

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.82	0.82	0.82	99
1	0.67	0.67	0.67	55
accuracy			0.77	154
macro avg	0.75	0.75	0.75	154
weighted avg	0.77	0.77	0.77	154



Random Forest Confusion Matrix:

```
[[81 18]
```

```
[18 37]]
```

Random Forest ROC-AUC Score: 0.84

Training SVM...

SVM Classification Report:

	precision	recall	f1-score	support
0	0.81	0.83	0.82	99
1	0.68	0.65	0.67	55
accuracy			0.77	154
macro avg	0.75	0.74	0.74	154
weighted avg	0.76	0.77	0.77	154

SVM Confusion Matrix:

```
[[82 17]
```

```
[19 36]]
```

SVM ROC-AUC Score: 0.83

Training XGBoost...

XGBoost Classification Report:


	precision	recall	f1-score	support
0	0.84	0.78	0.81	99
1	0.65	0.73	0.68	55
accuracy			0.76	154
macro avg	0.74	0.75	0.75	154
weighted avg	0.77	0.76	0.76	154

XGBoost Confusion Matrix:

```
[[77 22]
```

```
[15 40]]
```

XGBoost ROC-AUC Score: 0.81



Best Model: Random Forest with AUC = 0.84  
Best model saved as Random Forest\_pipeline.pkl